

Data

The data for this example consists of N=2,734 driver complaint to the National Transportation and Highway Safety Administration, NHTSA, regarding their automobile. It contains not only the complaints but also other information about the reason for their complaint and the type and condition of their automobile. The NHTSA uses these complaints to identify potential needs for automobile recalls.

The primary objective is to build a model that predicts whether the car was involved in a crash using the complaint and automobile characteristics.

Import Packages

In [1]:

```
# classes provided for the course
from Class_replace_impute_encode import ReplaceImputeEncode
from Class_regression import logreg
from sklearn.linear_model import LogisticRegression

from Class_tree import DecisionTree
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import export_graphviz
from pydotplus.graphviz import graph_from_dot_data
import graphviz

import pandas as pd
import numpy as np
import string
from nltk import pos_tag
from nltk.tokenize import word_tokenize
from nltk.stem.snowball import SnowballStemmer
from nltk.stem import WordNetLemmatizer
from nltk.corpus import wordnet as wn
from nltk.corpus import stopwords
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.decomposition import TruncatedSVD
```

The User Functions

my_analyzer(s) - Called by the sklearn Count & TFIDF Vectorizer's

The following functions will be used to customize the parse, pos, stop, stem process necessary for text analysis. These are done using the **NLTK** package, customized to remove certain words and symbols, and handle synonyms.

In [2]:

```
def my_analyzer(s):
    # Synonym List
    syns = {'veh': 'vehicle', 'car': 'vehicle', 'chev': 'chevrolet', \
            'chevy': 'chevrolet', 'air bag': 'airbag', \
            'seat belt': 'seatbelt', 'n't': 'not', 'to30': 'to 30', \
            'wont': 'would not', 'cant': 'can not', 'cannot': 'can not', \
            'couldnt': 'could not', 'shouldnt': 'should not', \
            'wouldnt': 'would not', 'straightforward': 'straight forward' }

    # Preprocess String s
    s = s.lower()
    # Replace special characters with spaces
    s = s.replace('-', ' ')
    s = s.replace('_', ' ')
    s = s.replace(',', '. ')
    # Replace not contraction with not
    s = s.replace("n't", " not")
    s = s.replace("n't", " not")
    # Tokenize
    tokens = word_tokenize(s)
    #tokens = [word.replace(',','') for word in tokens ]
    tokens = [word for word in tokens if ('*' not in word) and \
```

```

        ("'" != word) and ("``" != word) and \
        (word!='description') and (word !='dtype') \
        and (word != 'object') and (word!="'s")]

# Map synonyms
for i in range(len(tokens)):
    if tokens[i] in syns:
        tokens[i] = syns[tokens[i]]

# Remove stop words
punctuation = list(string.punctuation)+['..', '...']
pronouns = ['i', 'he', 'she', 'it', 'him', 'they', 'we', 'us', 'them']
others = ["'d", "co", "ed", "put", "say", "get", "can", "become", \
        "los", "sta", "la", "use", "iii", "else"]
stop = stopwords.words('english') + punctuation + pronouns + others
filtered_terms = [word for word in tokens if (word not in stop) and \
        (len(word)>1) and (not word.replace('.', '', 1).isnumeric()) \
        and (not word.replace("'", '', 2).isnumeric())]

# Lemmatization & Stemming - Stemming with WordNet POS
# Since lemmatization requires POS need to set POS
tagged_words = pos_tag(filtered_terms, lang='eng')
# Stemming with for terms without WordNet POS
stemmer = SnowballStemmer("english")
wn_tags = {'N':wn.NOUN, 'J':wn.ADJ, 'V':wn.VERB, 'R':wn.ADV}
wnl = WordNetLemmatizer()
stemmed_tokens = []
for tagged_token in tagged_words:
    term = tagged_token[0]
    pos = tagged_token[1]
    pos = pos[0]
    try:
        pos = wn_tags[pos]
        stemmed_tokens.append(wnl.lemmatize(term, pos=pos))
    except:
        stemmed_tokens.append(stemmer.stem(term))
return stemmed_tokens

def display_topics(lda, terms, n_terms=15):
    for topic_idx, topic in enumerate(lda):
        if topic_idx > 8:
            break
        message = "Topic #%d: " %(topic_idx+1)
        print(message)
        abs_topic = abs(topic)
        topic_terms_sorted = \
            [[terms[i], topic[i]] \
             for i in abs_topic.argsort()[::-n_terms - 1:-1]]

        k = 5
        n = int(n_terms/k)
        m = n_terms - k*n
        for j in range(n):
            l = k*j
            message = ''
            for i in range(k):
                if topic_terms_sorted[i+l][1]>0:
                    word = "+"+topic_terms_sorted[i+l][0]
                else:
                    word = "-"+topic_terms_sorted[i+l][0]
                message += '{:<15s}'.format(word)
            print(message)
        if m> 0:
            l = k*n
            message = ''

```

```

for i in range(m):
    if topic_terms_sorted[i+1][1]>0:
        word = "+" + topic_terms_sorted[i+1][0]
    else:
        word = "-" + topic_terms_sorted[i+1][0]
    message += '{:<15s}'.format(word)
print(message)
print("")
return

```

Attribute Map for Preprocessing Data

The following attribute map describes the data features. The attribute 'crashed' is the target in this problem. The attribute 'description' is a text screen that contains the driver's complaint.

Attributes with a **2** as the first number are nominal. There are two nominal attributes in the data, **Year**, **make**, **model** and attributes **T1-T9**. The topic attribute is the text topic cluster number. The attributes T1-T9 are the scores for individual documents for the topic cluster. In modeling, usually either the cluster topic assignment in the attribute **topic** or the topic scores in **T1-T9** will enter the model.

Attributes with a **3** as their first number are attributes that will be ignored and will not be encoded or returned in the encoded dataframe. In this case, the attribute **nhtsa_id** is ignored because they either have too many classes or have more than 50% missing.

In [3]:

```

attribute_map = {
    'nhtsa_id': [3, (0, 1e+12), [0, 0]],
    'Year': [2, (2003, 2004, 2005, 2006, 2007, 2008, 2009, 2010, 2011), [0, 0]],
    'make': [2, ('CHEVROLET', 'PONTIAC', 'SATURN'), [0, 0]],
    'model': [2, ('COBALT', 'G5', 'HHR', 'ION', 'SKY', 'SOLSTICE'), [0, 0]],
    'description': [3, (''), [0, 0]],
    'crashed': [1, ('N', 'Y'), [0, 0]],
    'abs': [1, ('N', 'Y'), [0, 0]],
    'mileage': [0, (1, 200000), [0, 0]],
    'topic': [2, (0, 1, 2, 3, 4, 5, 6, 7, 8), [0, 0]],
    'T1': [0, (-1e+8, 1e+8), [0, 0]],
    'T2': [0, (-1e+8, 1e+8), [0, 0]],
    'T3': [0, (-1e+8, 1e+8), [0, 0]],
    'T4': [0, (-1e+8, 1e+8), [0, 0]],
    'T5': [0, (-1e+8, 1e+8), [0, 0]],
    'T6': [0, (-1e+8, 1e+8), [0, 0]],
    'T7': [0, (-1e+8, 1e+8), [0, 0]],
    'T8': [0, (-1e+8, 1e+8), [0, 0]],
    'T9': [0, (-1e+8, 1e+8), [0, 0]]
}

```

Read the Data File

The following code reads the Excel data file using Pandas. The maximum column width in Pandas needs to be increased to ensure the text is read without truncation.

In [4]:

```

# Increase column width to let pandas read large text columns
pd.set_option('max_colwidth', 32000)
# Read NHTSA Comments
file_path = '/Users/Home/Desktop/python/Excel/'
df = pd.read_excel(file_path + "GMC_Complaints.xlsx")

```

Create Program Control Attributes

The **files** list is a list of the documents that will be processed. The remaining attributes are used to turn on and off tagging, stop words and stemming.

In [5]:

```

# Setup program constants
n_comments = len(df['description']) # Number of wine reviews
m_features = None # Number of SVD Vectors
s_words = 'english' # Stop Word Dictionary
comments = df['description'] # place all text reviews in reviews

```

```
n_topics = 9          # number of topic clusters to extract
max_iter = 10         # maximum number of iterations
max_df = 0.5          # learning offset for LDAmax proportion of docs/reviews allowed for a term
```

Tokenization, POS Tagging, Stop Removal & Stemming

There are two methods for text analysis: **TF-IDF** and **Counts**. The first is the term-frequency/inverse document frequency weighting. Raw term counts are weighted by the number of documents in which they appear. This down weights common terms used in all documents and up weights terms found in document clusters.

The second approach **Counts** is simply raw term frequencies. The term frequency matrix, **tf** contains the number of times a term appears in each review. The rows of this matrix are the reviews, $N=13,515$, and the columns are the terms. Most of the entries in this matrix are zero. That is not every term appears in every review. This matrix can be transformed using $\log(f+1)$ or binary.

The term descriptions are stored in a python list **terms** which has the same number of columns as **tf**, the maximum number of extracted terms. The term order is alphabetical rather than term frequency.

In [6]:

```
# Create Word Frequency by Review Matrix using Custom Analyzer
cv = CountVectorizer(max_df=0.95, min_df=2, max_features=m_features,\
                    analyzer=my_analyzer, ngram_range=(1,2))
tf = cv.fit_transform(comments)
terms = cv.get_feature_names()
term_sums = tf.sum(axis=0)
term_counts = []
for i in range(len(terms)):
    term_counts.append([terms[i], term_sums[0,i]])
def sortSecond(e):
    return e[1]
term_counts.sort(key=sortSecond, reverse=True)
print("\nTerms with Highest Frequency:")
for i in range(10):
    print('{:<15s}{:>5d}'.format(term_counts[i][0], term_counts[i][1]))
print("")
```

Terms with Highest Frequency:

| | |
|---------|------|
| vehicle | 6996 |
| steer | 2924 |
| contact | 2604 |
| power | 2131 |
| failure | 1745 |
| drive | 1670 |
| problem | 1466 |
| turn | 1256 |
| recall | 1239 |
| go | 1207 |

Create TFIDF Matrix

The following code create the TFIDF by transforming the term frequency matrix **tf**.

In [7]:

```
# Modify tf, term frequencies, to TF/IDF matrix from the data
print("Conducting Term/Frequency Matrix using TF-IDF")
tfidf_vect = TfidfTransformer(norm=None, use_idf=True) #set norm=None
tf = tfidf_vect.fit_transform(tf)

term_idf_sums = tf.sum(axis=0)
term_idf_scores = []
for i in range(len(terms)):
    term_idf_scores.append([terms[i], term_idf_sums[0,i]])
print("The Term/Frequency matrix has", tf.shape[0], " rows, and", \
      tf.shape[1], " columns.")
print("The Term list has", len(terms), " terms.")
term_idf_scores.sort(key=sortSecond, reverse=True)
print("\nTerms with Highest TF-IDF Scores:")
```

```
for i in range(10):
    print('{<15s}{:>8.2f}'.format(term_idf_scores[i][0], \
        term_idf_scores[i][1]))
```

Conducting Term/Frequency Matrix using TF-IDF

The Term/Frequency matrix has 2734 rows, and 3277 columns.

The Term list has 3277 terms.

Terms with Highest TF-IDF Scores:

| | |
|---------|---------|
| vehicle | 8162.39 |
| contact | 5371.28 |
| steer | 5114.56 |
| power | 4081.67 |
| failure | 3553.85 |
| problem | 3208.86 |
| drive | 2958.99 |
| recall | 2788.68 |
| turn | 2765.90 |
| go | 2757.42 |

Singular Value Decomposition

Use SVD to decompose the TFIDF matrix **tf**. This is called Latent Semantic Analysis - **LSA**.

In [8]:

```
# In sklearn, SVD is synonymous with LSA (Latent Semantic Analysis)
uv = TruncatedSVD(n_components=n_topics, algorithm='arpack', \
    tol=0, random_state=12345)

U = uv.fit_transform(tf)

# Display the topic selections
print("\n***** GENERATED TOPICS *****")
display_topics(uv.components_, terms, n_terms=15)
***** GENERATED TOPICS *****
```

Topic #1:

| | | | | |
|----------|--------|----------|----------|----------|
| +vehicle | +steer | +contact | +power | +problem |
| +would | +go | +recall | +failure | +drive |
| +turn | +time | +start | +gm | +dealer |

Topic #2:

| | | | | |
|----------|---------------|------------|--------|---------|
| -contact | -failure | -mileage | -state | -own |
| +problem | -manufacturer | +go | -fuel | -repair |
| -current | +start | -chevrolet | -mph | +fix |

Topic #3:

| | | | | |
|--------|--------|--------|-----------|---------|
| -steer | -power | +fuel | +ignition | +key |
| +start | +pump | -drive | +switch | +saturn |
| -go | +leak | -turn | -wheel | +smell |

Topic #4:

| | | | | |
|----------|---------|---------|---------|-------|
| -power | -steer | +brake | +front | +tire |
| +air | +side | -fuel | +bag | +door |
| +vehicle | +deploy | -recall | +driver | +hit |

Topic #5:

| | | | | |
|---------|-----------|---------|----------|-----------|
| +fuel | -ignition | -key | +pump | -start |
| -switch | +leak | +recall | +smell | +gasoline |
| -saturn | -would | +tank | -contact | +gas |

Topic #6:

| | | | | |
|----------|---------|----------|---------|------------|
| +door | +open | +handle | -brake | -start |
| -vehicle | +side | +issue | +insid | +key |
| +driver | -saturn | +plastic | +safety | +passenger |

Topic #7:

| | | | | |
|-----|------|------|-------|--------|
| -gm | -bag | -air | +fuel | +start |
|-----|------|------|-------|--------|

| | | | | |
|---------|---------|----------|----------|-------|
| +tire | -recall | +brake | +vehicle | -part |
| -deploy | +door | -problem | -crash | +turn |

Topic #8:

| | | | | |
|-----------|----------|---------|-----------|----------|
| +key | -start | -saturn | +ignition | -door |
| +remove | -problem | +gear | -ion | +release |
| +position | -cold | +gm | +turn | +shift |

Topic #9:

| | | | | |
|----------|------------|------------|--------|-------|
| +tire | -fuel | +firestone | +tread | -door |
| -vehicle | +rim | +new | +tell | -bag |
| -deploy | -passenger | +saturn | -seat | +part |

Add Topic Scores to Dataframe

The matrix **U** contains the SVD calculations that can be used to assign each document to a topic group. This code examines the SVD matrix **U** and assigns topic groups to each document. It then augments the original data frame with the document topic group assignments and **U**.

In [9]:

```
# Store topic group for each doc in topics[]
topics = [0] * n_comments
topic_counts = [0] * (n_topics+1)
for i in range(n_comments):
    max = abs(U[i][0])
    topics[i] = 0
    for j in range(n_topics):
        x = abs(U[i][j])
        if x > max:
            max = x
            topics[i] = j
    topic_counts[topics[i]] += 1

print('{:<6s}{:>8s}{:>8s}'.format("TOPIC", "COMMENTS", "PERCENT"))
for i in range(n_topics):
    print('{:>3d}{:>10d}{:>8.1%}'.format((i+1), topic_counts[i], \
        topic_counts[i]/n_comments))

# Create comment_scores[] and assign the topic groups
comment_scores = []
for i in range(n_comments):
    u = [0] * (n_topics+1)
    u[0] = topics[i]
    for j in range(n_topics):
        u[j+1] = U[i][j]
    comment_scores.append(u)

# Augment Dataframe with topic group information
cols = ["topic"]
for i in range(n_topics):
    s = "T"+str(i+1)
    cols.append(s)
df_topics = pd.DataFrame.from_records(comment_scores, columns=cols)
df = df.join(df_topics)
```

| TOPIC | COMMENTS | PERCENT |
|-------|----------|---------|
|-------|----------|---------|

| | | |
|---|------|-------|
| 1 | 1764 | 64.5% |
| 2 | 433 | 15.8% |
| 3 | 99 | 3.6% |
| 4 | 157 | 5.7% |
| 5 | 117 | 4.3% |
| 6 | 52 | 1.9% |

| | | |
|---|----|------|
| 7 | 37 | 1.4% |
| 8 | 49 | 1.8% |
| 9 | 26 | 1.0% |

Logistic Regression

In [10]:

```
target = 'crashed'
# Drop data with missing values for target (price)
drops= []
for i in range(df.shape[0]):
    if pd.isnull(df['crashed'][i]):
        drops.append(i)
df = df.drop(drops)
df = df.reset_index()

encoding = 'one-hot'
scale = None # Interval scaling: Use 'std', 'robust' or None
# drop=False - do not drop last category - used for Decision Trees
rie = ReplaceImputeEncode(data_map=attribute_map, nominal_encoding=encoding, \
                           interval_scale = scale, drop=True, display=True)
encoded_df = rie.fit_transform(df)
varlist = [target, 'T1', 'T2', 'T3', 'T4', 'T5', 'T6', 'T7', 'T8', 'T9']
X = encoded_df.drop(varlist, axis=1)
y = encoded_df[target]
np_y = np.ravel(y) #convert dataframe column to flat array
col = rie.col
for i in range(len(varlist)):
    col.remove(varlist[i])

lr = LogisticRegression(C=1e+16, tol=1e-16)
lr = lr.fit(X,np_y)

logreg.display_coef(lr, X.shape[1], 2, col)
logreg.display_binary_metrics(lr, X, y)
```

***** Data Preprocessing *****

Features Dictionary Contains:

10 Interval,
2 Binary, and
4 Nominal Attribute(s).

Data contains 2734 observations & 19 columns.

Attribute Counts

| | Missing | Outliers |
|---------------|---------|----------|
| nthsa_id..... | 0 | 0 |
| Year..... | 0 | 0 |
| make..... | 0 | 0 |
| model..... | 0 | 0 |
| description.. | 0 | 0 |
| crashed..... | 0 | 0 |
| abs..... | 18 | 0 |
| mileage..... | 419 | 15 |
| topic..... | 0 | 0 |
| T1..... | 0 | 0 |
| T2..... | 0 | 0 |
| T3..... | 0 | 0 |
| T4..... | 0 | 0 |
| T5..... | 0 | 0 |
| T6..... | 0 | 0 |
| T7..... | 0 | 0 |
| T8..... | 0 | 0 |
| T9..... | 0 | 0 |

Coefficients:

| | |
|-------------|---------|
| Intercept.. | -1.2659 |
| mileage.... | -0.0000 |
| abs..... | 0.4808 |
| Year0..... | -0.5004 |
| Year1..... | -0.4877 |
| Year2..... | -0.8197 |
| Year3..... | -0.6183 |
| Year4..... | -0.4913 |
| Year5..... | -0.2074 |
| Year6..... | -0.1843 |
| Year7..... | -0.5706 |
| make0..... | 0.0664 |
| make1..... | -0.1984 |
| model0..... | 0.3771 |
| model1..... | 0.8236 |
| model2..... | -0.3107 |
| model3..... | -0.4256 |
| model4..... | -0.7083 |
| topic0..... | 0.7417 |
| topic1..... | 1.0639 |
| topic2..... | -1.5045 |
| topic3..... | 3.6721 |
| topic4..... | -2.6114 |
| topic5..... | -0.5026 |
| topic6..... | 1.6533 |
| topic7..... | -1.7375 |

Model Metrics

| | |
|-----------------------------|--------|
| Observations..... | 2734 |
| Coefficients..... | 26 |
| DF Error..... | 2708 |
| Mean Absolute Error..... | 0.3247 |
| Avg Squared Error..... | 0.1622 |
| Accuracy..... | 0.7714 |
| Precision..... | 0.8011 |
| Recall (Sensitivity)..... | 0.1975 |
| F1-Score..... | 0.3169 |
| MISC (Misclassification)... | 22.9% |
| class 0..... | 1.8% |
| class 1..... | 80.2% |

Confusion

| Matrix | Class 0 | Class 1 |
|--------------|---------|---------|
| Class 0..... | 1964 | 36 |
| Class 1..... | 589 | 145 |

Decision Tree

In [11]:

```

scale = None # Not needed for Decision Trees
rie = ReplaceImputeEncode(data_map=attribute_map, nominal_encoding=encoding, \
                           interval_scale = scale, drop=False, display=True)
encoded_df = rie.fit_transform(df)
varlist = [target, 'T1', 'T2', 'T3', 'T4', 'T5', 'T6', 'T7', 'T8', 'T9']
X = encoded_df.drop(varlist, axis=1)
y = encoded_df[target] # These are not standardized
np_y = np.ravel(y) #convert dataframe column to flat array
col = rie.col
for i in range(len(varlist)):
    col.remove(varlist[i])

dtc = DecisionTreeClassifier(max_depth=7, \
                             min_samples_split=5, min_samples_leaf=5)

```



```
dtc = dtc.fit(X,np_y)
```

```
DecisionTree.display_importance(dtc, col, plot=False)
```

```
DecisionTree.display_binary_metrics(dtc, X, y)
```

```
***** Data Preprocessing *****
```

Features Dictionary Contains:

10 Interval,

2 Binary, and

4 Nominal Attribute(s).

Data contains 2734 observations & 19 columns.

Attribute Counts

| | Missing | Outliers |
|---------------|---------|----------|
| nthsa_id..... | 0 | 0 |
| Year..... | 0 | 0 |
| make..... | 0 | 0 |
| model..... | 0 | 0 |
| description.. | 0 | 0 |
| crashed..... | 0 | 0 |
| abs..... | 18 | 0 |
| mileage..... | 434 | 15 |
| topic..... | 0 | 0 |
| T1..... | 0 | 0 |
| T2..... | 0 | 0 |
| T3..... | 0 | 0 |
| T4..... | 0 | 0 |
| T5..... | 0 | 0 |
| T6..... | 0 | 0 |
| T7..... | 0 | 0 |
| T8..... | 0 | 0 |
| T9..... | 0 | 0 |

FEATURE.... IMPORTANCE

| | |
|-------------|--------|
| topic3..... | 0.4967 |
| mileage.... | 0.1480 |
| make2..... | 0.0899 |
| model2..... | 0.0408 |
| topic4..... | 0.0364 |
| topic2..... | 0.0328 |
| topic6..... | 0.0279 |
| model5..... | 0.0257 |
| Year4..... | 0.0215 |
| topic1..... | 0.0192 |
| Year2..... | 0.0179 |
| topic0..... | 0.0131 |
| Year3..... | 0.0095 |
| abs..... | 0.0090 |
| make0..... | 0.0062 |
| model0..... | 0.0042 |
| Year1..... | 0.0012 |
| Year0..... | 0.0000 |
| Year5..... | 0.0000 |
| Year6..... | 0.0000 |
| Year7..... | 0.0000 |
| Year8..... | 0.0000 |
| make1..... | 0.0000 |
| model1..... | 0.0000 |
| model3..... | 0.0000 |
| model4..... | 0.0000 |
| topic5..... | 0.0000 |
| topic7..... | 0.0000 |
| topic8..... | 0.0000 |

Model Metrics

| | |
|-----------------------------|--------|
| Observations..... | 2734 |
| Features..... | 29 |
| Maximum Tree Depth..... | 7 |
| Minimum Leaf Size..... | 5 |
| Minimum split Size..... | 5 |
| Mean Absolute Error..... | 0.3107 |
| Avg Squared Error..... | 0.1553 |
| Accuracy..... | 0.7835 |
| Precision..... | 0.8142 |
| Recall (Sensitivity)..... | 0.2507 |
| F1-Score..... | 0.3833 |
| MISC (Misclassification)... | 21.7% |
| class 0..... | 2.1% |
| class 1..... | 74.9% |

Confusion

| Matrix | Class 0 | Class 1 |
|--------------|---------|---------|
| Class 0..... | 1958 | 42 |
| Class 1..... | 550 | 184 |